

Project 4: Reinforcement Learning

Train a Smartcab to Drive

The setting

In the project “Train a Smartcab to Drive”, a smartcab operating in an idealized grid-like city is given. There are traffic lights at each intersection and other cars present. The smartcab gets a reward for obeying traffic rules and a penalty for not obeying traffic rules or causing an accident. Goal of this project is to implement a learning agent for the smartcab that should learn an optimal policy for driving on city roads, obeying traffic rules correctly, and trying to reach the destination within a goal time based on the rewards and penalties it gets.

Step 1 - Implement a basic driving agent

In the first step, I let the smart cab take a random action out of the possibilities of:

- doing nothing (state ‘None’)
- driving forward (state ‘forward’)
- turning left (state ‘left’)
- turning right (state ‘right’)

In this step, the smartcab does not learn from the results of it's actions and has unlimited time to reach the goal.

I put the printed output of my test run into 'output_first_text.txt'.

```
In [279]: def read_in_my_text(file_name, print_lines):  
          with open (file_name, "r") as myfile:  
              output = myfile.read()  
              count_reached = output.count('Primary agent has reached destina  
tion!')  
              count_aborted = output.count('Trial aborted.')  
              lines = output.split('\n')  
              print "Rate of reaching the destination {}".format(count_reache  
d*1.0/(count_reached + count_aborted))  
              if print_lines == True:  
                  print "Last 3 lines:"  
                  for line in lines[-3:]:  
                      print line
```

```
In [280]: read_in_my_text("outputs/out_1_1.txt", False)  
  
Rate of reaching the destination 0.19801980198
```

```
In [281]: read_in_my_text("outputs/out_1_2.txt", False)  
  
Rate of reaching the destination 0.25
```

```
In [282]: read_in_my_text("outputs/out_1_3.txt", False)  
  
Rate of reaching the destination 0.2
```

```
In [283]: read_in_my_text("outputs/out_1_4.txt", False)  
  
Rate of reaching the destination 0.23
```

```
In [284]: read_in_my_text("outputs/out_1_5.txt", False)  
  
Rate of reaching the destination 0.16
```

The smartcab reaches the destination in time in around 20% of the trials. If the smartcab reaches the destination is a question of chance. The actions are chosen randomly, independent of the destination and the received rewards.

Step 2 - Identify and update state

For the states, I chose a tuple of 'next_waypoint', 'light', 'oncoming' and 'left'. I didn't take 'right' into account, because no matter what our smartcab is going to do, the traffic from the right side has no influence on it.

Step 3 - Implement Q-Learning

In this step I implemented the q-learning algorithm with a learning rate of 0.5 and a discount factor of 0.5. The next action is always chosen based on the best estimate based on the current state. I didn't use an exploration rate.

```
In [285]: read_in_my_text("outputs/out_3_1.txt", True)
```

```
Rate of reaching the destination 0.911764705882
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [286]: read_in_my_text("outputs/out_3_2.txt", True)
```

```
Rate of reaching the destination 0.93137254902
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [287]: read_in_my_text("outputs/out_3_3.txt", True)
```

```
Rate of reaching the destination 0.06
Last 3 lines:
Environment.step(): Primary agent ran out of time! Trial aborted.
Environment.step(): Primary agent ran out of time! Trial aborted.
Environment.step(): Primary agent ran out of time! Trial aborted.
```

```
In [288]: read_in_my_text("outputs/out_3_4.txt", True)
```

```
Rate of reaching the destination 0.970297029703
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [289]: read_in_my_text("outputs/out_3_5.txt", True)
```

```
Rate of reaching the destination 0.950980392157
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

Interestingly, the smartcab reaches the destination quite often in time. There seem to be some cases, when the smartcab gets stuck and doesn't reach the destination in time.

Step 4 - Enhance the driving agent

I played around with the learning rate and calculated it using the time step. I stayed with a discount factor of 0.1. I tried several exploration rates, but that lead to not reaching the goal more often.

```
In [290]: read_in_my_text("outputs/out_5_1.txt", True)
```

```
Rate of reaching the destination 0.950495049505
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [291]: read_in_my_text("outputs/out_5_2.txt", True)
```

```
Rate of reaching the destination 0.98
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [292]: read_in_my_text("outputs/out_5_3.txt", True)
```

```
Rate of reaching the destination 0.94
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.step(): Primary agent ran out of time! Trial aborted.
```

```
In [293]: read_in_my_text("outputs/out_5_4.txt", True)
```

```
Rate of reaching the destination 0.98
Last 3 lines:
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
Environment.act(): Primary agent has reached destination!
```

```
In [294]: read_in_my_text("outputs/out_5_5.txt", True)
```

```
Rate of reaching the destination 0.92
```

```
Last 3 lines:
```

```
Environment.act(): Primary agent has reached destination!
```

```
Environment.act(): Primary agent has reached destination!
```

```
Environment.act(): Primary agent has reached destination!
```

In addition to the changes I made to the learning rate, I changed the initialization from my q table from initializing it with zeros to using random values between 0 and 4. With my first attempt the smartcab got stuck most of the time.

My final agent finds a close to optimal policy. The rewards increase with time and the smartcab reaches the destination in time most of the time.